



Berner  
Fachhochschule

CAS Java Microservice Development

# Java Persistence API - O/R-Mapping

Simon Martinelli, [simon.martinelli@bfh.ch](mailto:simon.martinelli@bfh.ch), v2025.10

# Content

1. Objection-Relational-Mapping
2. Relationship Mapping
3. Lazy Loading

# O/R-Mapping

# Access Type - Field Access

```
@Entity  
public class Employee {  
  
    @Id  
    private Integer id;  
  
    public Integer getId() {  
        return id;  
    }  
  
    public Integer setId() {  
        this.id = id;  
    }  
}
```

# Access Type - Property Access

```
@Entity  
public class Employee {  
  
    private Integer id;  
  
    @Id  
    public Integer getId() {  
        return id;  
    }  
  
    public Integer setId() {  
        this.id = id;  
    }  
}
```

# Access Type Options

- Different access types possible per class
- Mixing of access types in an inheritance hierarchy

```
@Entity  
@Access(FIELD)  
public class Vehicle {  
  
    @Transient  
    double fuelEfficiency;  
  
    @Access(PROPERTY)  
    protected double getDbFuelEfficiency() {  
        return convertToImperial(fuelEfficiency);  
    }  
}
```

# Mapping

```
@Entity  
@Table(name = "employees")  
public class Employee {  
  
    @Id  
    @Column(name = "employees_id")  
    private Integer id;  
  
    @Basic(optional = false)  
    private String name;  
}
```

# Persistent Data Types

- All primitive types, `String` and all wrapper classes (e.g., `Integer`), `BigDecimal`, Temporal Types
- `byte[]`, `Byte[]`, `char[]`, `Character[]`
- Enumerations
- References to
  - any other entity classes
  - Collections of Entities and wrapper classes and String, which are declared as `Collection`, `List`, `Set` or `Map`

# Java/SQL Type Mapping

- Defined implicitly by JDBC "Data Type Conversion Table"

[https://download.oracle.com/otn-pub/jcp/jdbc-4\\_1-mrel-spec/jdbc4.1-fr-spec.pdf?  
AuthParam=1572082201\\_b508cd35046b6f60a93969dd4f9cdce](https://download.oracle.com/otn-pub/jcp/jdbc-4_1-mrel-spec/jdbc4.1-fr-spec.pdf?AuthParam=1572082201_b508cd35046b6f60a93969dd4f9cdce)

- Explicitly by `@Column` annotation, e.g.

```
@Column(columnDefinition = "VARCHAR(20)"  
private String name;
```

- Product specific
  - JPA implementation
  - JDBC Driver

# Large Objects

- Storing data in BLOB or CLOB

```
public class Employee {  
  
    @Lob // BLOB  
    private byte[] picture;  
  
    @Lob // CLOB  
    private char[] largeText;  
  
}
```

# Enumerations

- Enumerations can be persisted. Either as ordinal value (position) or as a string (name of the constant)

```
public enum Color { RED, BLUE, GREEN}
```

```
@Enumerated(EnumType.ORDINAL) // Default  
private Color color;
```

```
@Enumerated(EnumType.STRING)  
private Color color;
```

Be careful when changing Enums!

# Temporal Types

Java 8 Date/Time API is recommended

- Without timezone
  - `java.time.LocalDate`
  - `java.time.LocalTime`
  - `java.time.LocalDateTime`
- With timezone
  - `java.time.OffsetTime`
  - `java.time.OffsetDateTime`

# Temporal Types Legacy Support

- Supported Types
  - `java.sql.Date` , `java.sql.Time` , `java.sql.Timestamp` ,
  - `java.util.Date`
  - `java.util.Calendar`
- `java.sql` types do not require further definition but for `java.util` the JDBC type must be specified
  - `TemporalType.DATE` , `TemporalType.TIME` , `TemporalType.TIMESTAMP`

```
@Temporal(TemporalType.DATE)
private java.util.Date dateOfBirth;
```

# Transient Attributes

- Attributes can be excluded from the persistence
- Either transiently using the keyword `transient`

```
transient private String translatedName;
```

- or if the attribute must be serializable using an annotation

```
@Transient  
private String translatedName;
```

# Entity Identity - The Primary Key

- Each entity class must have a primary attribute annotated with `@Id`
- An Id can be of the following types
  - Primitive Java types: `byte` , `short` , `int` , `long` , `char`
  - Wrapper classes: `Byte` , `Short` , `Integer` , `Long` , `Character`
  - Array of primitive types or wrapper classes
  - `java.lang.String` , `java.math.BigInteger`
  - Temporal types
  - (Floating Point types are also allowed, but are not recommended because of possible rounding errors)

# Primary Key Generation

- Primary Keys can be generated by JPA with the database
- There are four strategies: Sequence, Identity, UUID, Table, and Auto

| Source: <https://vladmihalcea.com/jpa-entity-identifier-sequence/>

# Generation Type

- SEQUENCE allows using a database sequence object to generate identifier values.  
This is the best generation strategy when using JPA and Hibernate.
- IDENTITY allows using a table identity column, like the MySQL AUTO\_INCREMENT.
- UUID generates a UUID as primary key.
- TABLE emulates the database sequence generator using a separate table. This is a terrible strategy, and you shouldn't use it.
- AUTO picks any of the previous strategies based on the underlying database capabilities. It shouldn't be used.

# Sequence

- For JPA and Hibernate, you should prefer using SEQUENCE if the relational database supports it because Hibernate cannot use automatic JDBC batching when persisting entities using the IDENTITY generator.
- Plus IDENTITY must execute the insert statement to get the ID

```
@Entity public class Employee {  
    @Id  
    @GeneratedValue(strategy = GenerationType.SEQUENCE)  
    public Integer id;  
}
```

# Sequence Generator

- By default, Hibernate creates a HIBERNATE\_SEQUENCE sequence
- If you want to use your own sequence, most often per table, you can use the `@SequenceGenerator` annotation

```
@Entity public class Employee {  
  
    @Id  
    @GeneratedValue(strategy = GenerationType.SEQUENCE, generator="employee_seq")  
    @SequenceGenerator(name = "employee_seq", sequenceName = "employee_seq")  
    public Integer id;  
  
}
```

Make sure that the allocation size matches the sequence allocation size!

# UUID Issues

- The index pages will be sparsely populated because each new UUID will be added randomly across the B+Tree clustered index.
- There are going to be more page splits because of the randomness of the Primary Key values
- The UUID is huge, needing twice as many bytes as a bigint column. Not only it affects the Primary Key but all the associated Foreign Keys as well.
- More, if you're using SQL Server, MySQL, or MariaDB, the default table is going to be organized as a Clustered Index, making all these problems even worse.
- So, you are better off avoiding using the UUID for entity identifiers. If you really need to generate unique identifiers from the application, then you are better off using a 64-bit time-sorted random TSID instead.



# Relationship Mapping

# Relationships

- Relationships between entities are principally represented by corresponding references or collections in the entity classes.
- They must be declared, and details are often necessary for O/R mapping and behavior.
- The following relationship characteristics play a role:
  - Direction
  - Cardinality
  - Aggregation, Composition

# Characteristics



- Unidirectional
- Bidirectional
- OneToOne
- ManyToOne
- OneToMany
- ManyToMany



# Owning and Inverse Side



- JPA distinguishes between **owning** and **inverse side**
- **The owning side is responsible for managing the relationship** in the database (Foreign Key)
- The **inverse side** is marked by the attribute `mappedBy`
- In unidirectional relationships, the inverse side is missing

# OneToOne, Unidirectional





```
public class Employee {  
    @OneToOne  
    private Address address;  
}
```

# ManyToOne, Unidirectional



```
public class Employee {  
    @ManyToOne  
    private Department department;  
}
```



# OneToMany, Bidirectional



```
public class Employee {  
    @OneToMany(mappedBy = "employee")  
    private Set<Phone> phones;  
}
```



```
public class Phone {  
    @ManyToOne(optional = false)  
    private Employee employee;  
}
```

# OneToMany, Unidirectional



```
public class Employee {  
    @OneToMany  
    @JoinColumn(name = "employee_id", nullable = false)  
    private Set<Phone> phones;  
}
```

# ManyToMany, Bidirectional




# ManyToMany, Bidirectional

```
public class Employee {  
  
    @ManyToMany(mappedBy = "employees")  
    private Set <Project> projects;  
}
```

```
public class Projects {  
  
    @ManyToMany  
    private Set<Employee> employees;  
}
```

# Embeddables

- A way to implement composition
- Embedded objects do not have their own identity
- Entity and embeddables are stored in the same table



EMPLOYEE	
PK	ID
	NAME
	SALARY
	STREET
	CITY
	STATE
	ZIP

# Example

```
@Embeddable  
public class Address {  
    private String street;  
    private String city;  
    private String state;  
    private String zip;  
}
```

```
@Entity  
public class Employee {  
    @Id  
    private int id;  
    private String name;  
    private long salary;  
    @Embedded  
    private Address address;  
}
```

# Collections of Non-Entities and Embeddables

```
@Entity  
public class Employee {  
  
    @ElementCollection  
    @Column(name = "PHONE_NUMBER")  
    private List<String> phoneNumbers;  
}
```

# Collections

- `java.util.Set` : unique according to `equals()`

```
@OneToMany
```

```
private Set<Phone> phones;
```

- `java.util.List` : ordered, can be sorted

```
@OneToMany @OrderBy("phonenumbers ASC")
```

```
private List<Phone> phones;
```

- `java.util.Map` : key/value pairs

```
@OneToMany @MapKey(name = "phonenumbers")
```

```
private Map<String, Phone> phones;
```

# Persistent Ordering

- The order of a list can be persistent

```
@Entity  
public class Employee {  
  
    @OneToMany  
    @OrderColumn(name = "PHONE_POS")  
    List<Phone> phones;  
}
```

# Enhanced Map Support

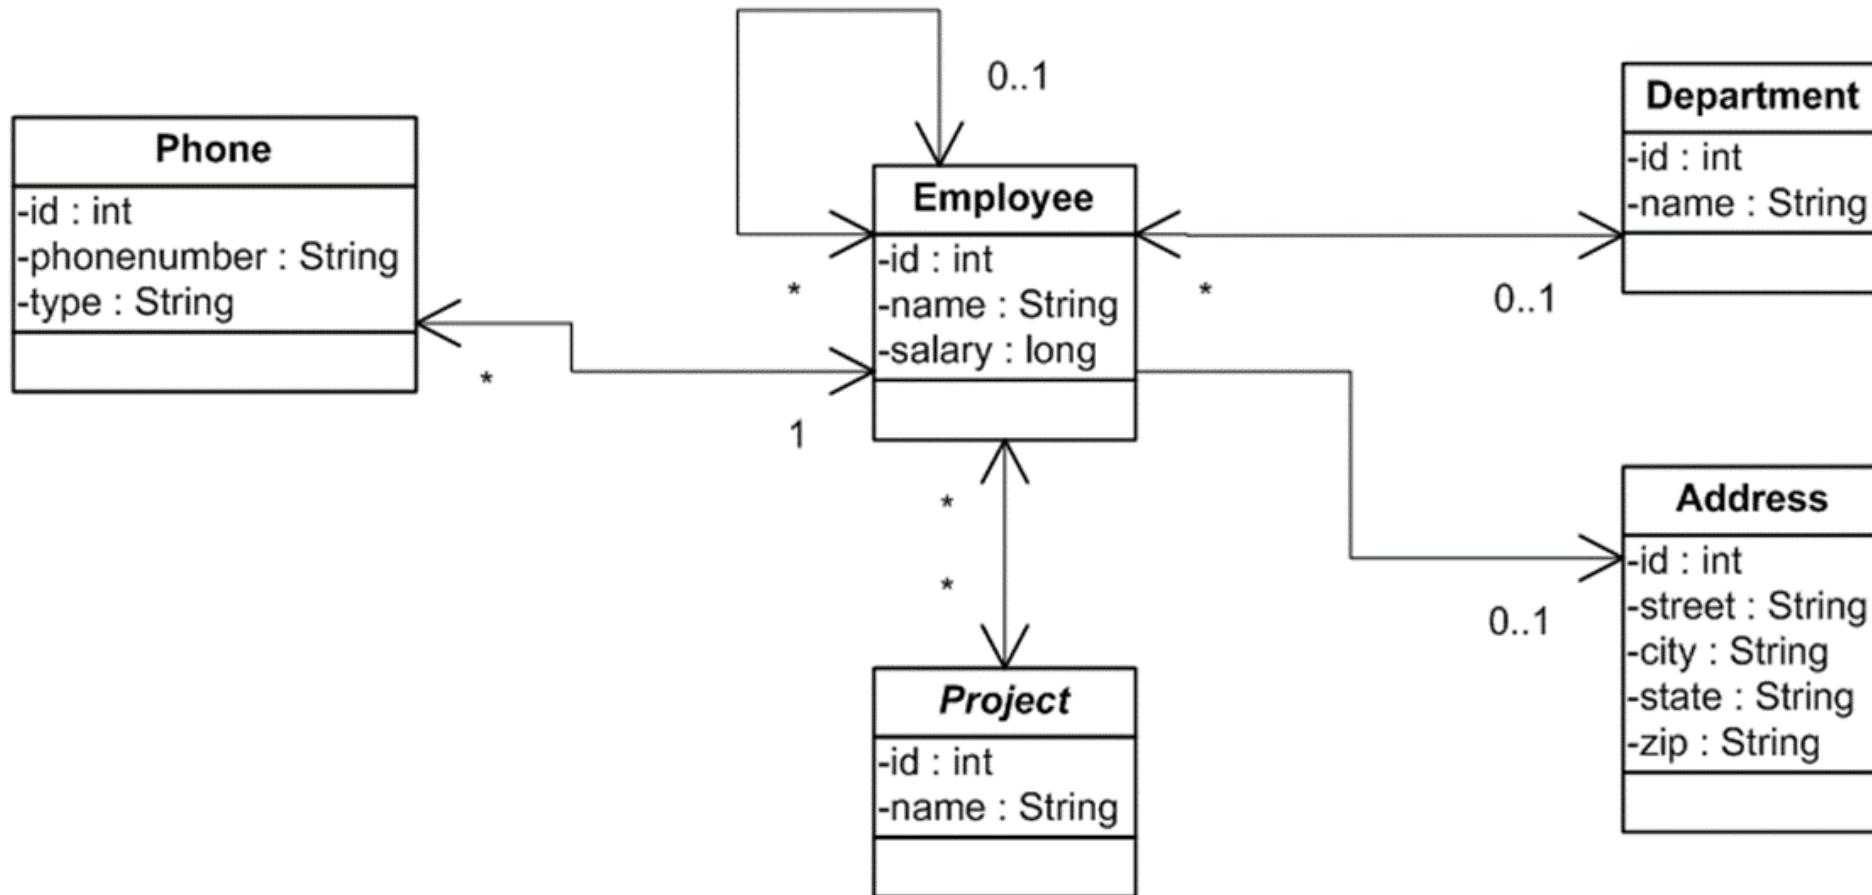
- Use objects, embeddables and entities as map key and value

```
@Entity  
public class Vehicle {  
  
    @OneToMany  
    @MapKeyJoinColumn(name="PART_ID")  
    Map<Part, Supplier> suppliers;  
}
```

# Exercise: Ids and Relationships

0. Remove Flyway and use Hibernate to create the database model
1. Implement the class model  
*(please use Integer instead of int)*
2. Select the Id generation strategy
3. Define the relationship types
4. Define the relationship mappings
5. Create Repositories
6. Test your work by creating n Employee with all relationships

# Entity Model




# Lazy Loading

# Lazy Loading

- Default with `OneToOne` and `ManyToOne`  
`FetchType.EAGER`
- Default with `OneToMany` and `ManyToMany`  
`FetchType.LAZY`
- The default behavior can be overridden

```
@OneToMany(fetch = FetchType.EAGER)  
private Set<Phone> phones;
```

# Problems with Lazy Loading



# Solutions

# Deactivate Lazy Loading

- Advantage
  - Simple `FetchType.EAGER`
- Disadvantages
  - May lead to performance issues because it may not JOIN the tables
  - Generates n+1 bombs
  - Hibernate can only eagerly load more than one List per Entity
- Recommendation
  - Don't use! Instead, set all `toOne`, and `ToMany` annotations to `FetchType.LAZY`

# Open Session In View Pattern

- The OpenSessionInView pattern is a pattern that assumes that in an application, the session (EntityManager) is not closed until the request is finished
- Advantage
  - The data is reloaded as needed
- Disadvantages
  - No clean separation of the layers
  - Not all data is loaded in the same transaction
- Recommendation: Don't use. This is an anti-pattern! [Article from Vlad](#)  
Caution: Spring Boot version  $\geq 2$  has activated this pattern by default  
`spring.jpa.open-in-view=true`

# JOIN FETCH

- Advantage

- Simple

```
select e from Employee e join fetch e.phones p order by e.name
```

- Disadvantage

- Leads to many queries

- Recommendation

- Use. Very useful and pretty clear what happens.

# JPA Entity Graph

```
@Entity  
@NamedEntityGraph(name="includePhones", attributeNodes={@NamedAttributeNode("phones")})  
public class Employee {  
}
```

```
EntityGraph includePhones = em.getEntityGraph("includePhones");  
Query query = em.createQuery("select e from Employee e order by e.name");  
query.setHint("javax.persistence.loadgraph", includePhones);  
return query.getResultList();
```

# JPA Entity Graph

- Advantage
  - EntityGraph is reusable
- Disadvantage
  - Leads to lots of annotations
- Recommendation
  - Use if you need to reuse the entity graphs